



A-OSP

The open format for proof of AI work

AI produces fluent output, not evidence. A-OSP turns AI-assisted work into typed, local, reconstructable proof — starting where regulated work already pays for it.

Formation round: €1.2M · 24 months · capital released against verifiable gates

Gianluca Conte · contegianluca@hotmail.com · Milan · 2026

draft · third-party names are category examples, no relationship or endorsement implied · not legal or financial advice

AI creates output. A-OSP creates reconstructable proof.

The investable point is not a better dashboard. **It is owning the evidence object — the open format of proof that every other tool will have to read. We are not selling a GRC suite; we are selling the substrate underneath it.**

The shift

AI makes producing work faster — and reconstructing proof and accountability more expensive

The object

Work becomes typed, traceable, reusable evidence: source, lineage, verification, declared gaps

This round


We are not funding the vision. We are funding the minimum commercial proof: paid pilots, conversion, cost-to-serve, channel

If the gates pass, A-OSP becomes the proof standard every tool, auditor and regulator must read — a category, not a feature. If they don't, it remains a technical, standardizable asset with contained downside.

THE PROBLEM

Boards buy defensible reconstructability, not «smart AI»

The auditor's question is never «is the model good?» — it is «who decided, on what sources, with what limits?»



In regulated work the problem is not speed — it is whether the work is admissible at all.

One segment first: multi-regime SMEs, via their advisors

Not «SME + enterprise + PA + model makers». One door, where the pain is already paid for.

Who

- Structured Italian SMEs already exposed to ≥ 2 regimes (231 + GDPR + NIS2/H&S)
- They already pay consultants for it — budget exists, buyer exists, urgency exists
- Reachable through the founder's own network: associations, OdV, advisory firms

Why this door

- Short sales cycle: bought by a compliance lead, no board resolution needed
- Each SME validates that one evidence object serves many regimes
- Volume produces references, cost-to-serve data and channel proof

Larger markets exist — but they are EXPANSION, not the ask (see appendix):

Regulated EU enterprise (finance, health, energy)

Digital government trust / sovereign


Model-maker & infra partnerships

The case base reaches them only after the SME gates pass. The round does not depend on them.

THE USE CASE

One evidence object, many lenses. Audit pack writes itself.

TODAY — same evidence, rebuilt 4 times



WITH A-OSP — one typed object, four lenses



Output: an audit pack ready for inspection in minutes — every claim reconstructable to source, what's missing declared.

The principle that scales: each new regime is a lens on the same object, not a new silo. A-OSP does not decide legal truth — it builds the inspectable object the decision rests on.

Two clocks force the spend — independent of AI hype

The regulatory clock

- AI Act in force (high-risk Dec 2027); fines up to €35M / 7% of turnover
- DORA applicable since 17 Jan 2025; NIS2 across 18 sectors; GDPR already baseline
- Proof is becoming a legal obligation, not a best practice

The operational clock

- AI is entering advisory, audit, policy and risk work right now
- Output grows faster than verifiability — the reconstruction gap widens monthly
- Inference cost falls; trust cost stays high — that spread is the margin

The trigger we sell into: an SME facing a 231 update, a NIS2 deadline, a DORA supply-chain request or a first AI-assisted deliverable under audit. The pain has a date. Public-trust / sovereignty tailwinds (eIDAS 2.0, EUDI) reinforce the direction but are not part of this round's sales motion — see appendix.

The prize is a market, not a niche — and it is consolidating now

The market behind the wedge

- **GRC software ~\$51B (2023)**; >\$140B/yr spent on compliance & security reviews
- **AI governance ~\$0.75B (2024) → ~\$5.6B by 2030**, ~40% CAGR (analyst range)
- Gartner: AI rules ×4 by 2030, touching 75% of economies; governance platforms pass \$1B


Comparables prove the prize

- **Vanta ~\$300M ARR, ~\$4.15B valuation** — now repositioning as a system of record for proving trust
- **Drata ~\$100M ARR, ~\$2B valuation** — the category sustains multi-billion outcomes
- Incumbents reach proof from automation; A-OSP reaches it from a neutral, open proof substrate they cannot retro-own

The wedge is not the ceiling: multi-regime SMEs are the beachhead into the proof layer of all AI-assisted regulated work. If A-OSP becomes the format others read, the return logic is category ownership, not a services line. Market figures are external analyst estimates, orders of magnitude to verify in DD — not A-OSP revenue.

THE OFFER

Sell paid proof first, platform later: assessment → subscription



The assessment is the wedge: it pays for itself, lands a reference, and produces the cost-to-serve and conversion data the round must prove. Platform value comes after — not before — paid pull.

What we charge, what it costs to serve, what we must prove

Every figure below is a HYPOTHESIS to validate at the first paid assessments — not a result.

Metric	Working assumption	How we validate it
Assessment price	€12–20k one-off	first 3 paid assessments to the founder's network
Subscription (SME)	€4–6k / year	conversion of assessment clients to maintenance
Cost to serve / seat	€70–250 (HYP)	measured per project: inference + cloud + support
Delivery cost / assessment	8–15 senior-days (HYP)	actuals logged on each engagement
Gross margin target	75–85% (HYP)	gross profit per workflow, once delivery is standardized
Channel share	15–20%	association / advisor agreements

Why the margin can hold: the substrate carries the epistemic load, so inference is non-frontier and small-context; no training capex; the channel keeps SME CAC near zero. None of this is asserted — it is what the first engagements are designed to measure.

From senior-days to software: how delivery scales

The honest risk: high-quality consulting that doesn't scale. Here is how A-OSP becomes a platform.

Year 1 — founder-led

- ~8–15 senior-days / assessment (HYP)
- ~80% founder-dependent
- mostly custom: learning the repeatable core
- throughput $\approx$ 8–12 assessments

Year 2 — templated

- ~4–8 senior-days / assessment
- ~40% founder; engineers + 1 GRC hire deliver
- 60% repeatable templates, 40% custom
- throughput $\approx$ 30–40 + first cohorts

Year 3 — channel + product

- ~1–3 days config / cohort onboarding
- <15% founder; channel delivers most
- software-like: subscription > services revenue
- throughput limited by channel, not by us

The inflection: the metric we watch is senior-days per €1 of recurring revenue. When it crosses below the software threshold — recurring revenue grows while delivery days stay flat — A-OSP has stopped being consulting. Reaching that point is exactly what G2–G3 measure.

The only thing that matters first: do they pay, and renew?

Next 90 days

- Entities incorporated, IP assigned, lead engineer identified
- 1–2 paid assessments signed from the network
- Pricing tested; rejections logged with reasons

Gate G0–G1

By 12 months

- ≥5 paid assessments delivered, actuals logged
- ≥1 channel (association / advisor) with a pilot cohort
- First conversions to maintenance; cost-to-serve measured

Gate G2 — paid-pull proof


By 24 months

- SME lens in GA; 25 SMEs + 4 mid-market
- Conversion ≥15%; NRR signal; channel repeatable
- Data room full enough to raise round 2

Gate G3 — repeatability


Falsification test: if assessments don't sell to the founder's own network, at this price, the commercial thesis dies early and cheap. That is a feature, not a risk we hide.

The pain already has a vendor: advisors, associations, OdV



Without the channel the bottom-up motion does not work — stated plainly.

What exists, what's missing, what the round funds



Real today (verifiable in data room)

- 4-layer architecture and first proof stage well advanced: verification, receipts, controlled writes, negative tests, fail-closed UI
- Built by one founder orchestrating AI agents with a documented method — technical direction demonstrated, not claimed

What the round closes (mapped backlog)

- End-to-end proof chain and audit pack; SME multi-regime lens in GA; cohort onboarding
- Reproducible deployment; security baseline / SOC 2 started; measured pilots

Provider-agnostic by design: models (frontier or local) are replaceable processors; truth stays inside the client's perimeter.

Open standard + company is the moat — not an org chart

It answers three questions at once: why the market trusts it, why incumbents can't own the format, why investors and grants coexist.

Why the market trusts it

A proof format no single vendor controls is one an auditor, a DPO or a procurement office can rely on. Inspectable beats proprietary where evidence is the product.

Why incumbents can't own it

An incumbent can copy features into its suite. It cannot retroactively own an open, neutral format already adopted by associations and held by a foundation.

Why capital types coexist

The non-profit attracts grants and public money; the OpCo attracts equity. Each funds what it should — neither contaminates the other's incentives.

The standard ambition belongs here, not in the revenue line. A-OSP does not need to sell to governments to be valuable — but to become respectable, neutral and adoptable as a format, the open core and a credible standards body are the mechanism. That is why the legal architecture is strategy, not paperwork. (Where it sells — PA, sovereign, model makers — stays in the appendix; why it can become a standard stays here.)

Who owns what. How the OpCo makes money if the core is open

Foundation owns

- The open specification of the evidence format (the standard)
- The trademark and the certification mark («A-OSP-compliant»)
- Open-core reference implementation, under a network-copyleft licence
- Neutral governance of the spec via a standards committee

OpCo owns

- All product IP: enterprise modules, managed runtime, integrations
- A durable commercial licence to the core for paid services and deployments
- The brand for commercial offerings; customer contracts and data-processing
- Equity, ESOP, the cap table — what investors actually buy

How the OpCo captures value while the core stays open

- Open core = trust and distribution; the things buyers actually pay for are NOT in the commons: SLAs, managed hosting, enterprise lenses, certifications, sovereign deployments, support
- The certification mark is the toll: «A-OSP-compliant» is granted by the Foundation but operationalized through OpCo tooling and services
- Investor neutrality concern resolved: the Foundation governs the spec (independent board); the OpCo cannot unilaterally fork the standard — which is precisely what makes the standard credible and therefore valuable


Every capital type knows what puts in, gets out, and is measured on

Funder	Puts in	Gets back	Return KPI
Angel / pre-seed	equity (€25–50k+)	OpCo upside, gated downside	paid pull · ARR · round 2 unlocked
Association	channel + reputation	revenue share + standards seat	members onboarded
Professional / advisor	clients + delivery	more value per hour + retention	revenue per senior-day
SME	assessment / subscription fee	less reconstruction + audit pack	hours saved per audit
Public / grant	non-dilutive capital	open standard + sovereignty	adoption · training reach
Corporate sponsor	pilot / partnership	provider-agnostic proof layer	integration · risk reduction
Model maker	adapter / co-sell	access to regulated adoption	regulated deployments

The point: the commons monetizes trust; the OpCo monetizes delivery, SLAs and integration; the channel monetizes distribution. Not «everyone buys SaaS» — each euro enters for a different, measurable reason. This round is only the first row; the others compound as the standard spreads.

A1–A2 is what we underwrite; A3–A5 is gated upside (not yet a promise)

ARR end-of-year (€M)



Category-capture scenario (HYPOTHESIS · gated, round 2+): if the standard is adopted, A5+ trends toward €50–100M+ ARR — Vanta (~\$300M ARR) and Drata (~\$100M ARR) prove the category supports it. Exit: GRC suites / model-makers buy a proof substrate (cf. Drata–SafeBase, \$250M). See THE CATEGORY + Appendix B.

Full P&L A0–A5, microfoundation formulas and sensitivity: in appendix.

A1–A2

what €1.2M underwrites: paid pull, conversion, cost-to-serve — the provable part

≈ €1.15M

cumulative net burn over the 24 funded months ≤ the round: the arithmetic closes

75–85%
(HYP)

gross-margin target, measured per workflow in pilots — not asserted

Grants

public capital always reported separately from ARR — funds product and commons, never revenue

A domain founder, plus the technical core the round hires

Gianluca Conte — founder & CEO

- Head of IT Audit & Compliance (Bocconi/SDA); ex Head of Internal Audit; 9 years secretary of a 231 supervisory body
- M.Sc. Cybersecurity, M.Sc. Auditing, applied mathematics; Pg.Dip. philosophy applied to AI
- IAPP AI Governance Europe, NIS2-CODAU and AODV231 think tanks, Clusit
- Built the prototype by orchestrating AI agents — technical direction demonstrated, not claimed

What the round hires

- Lead engineer / CTO: first name, condition precedent to closing (cash + equity)
- Core of 3–5 engineers (proof runtime, platform, security) on the mapped backlog
- GRC lead and customer success mid-way; fractional CFO and counsel
- Advisory board of 3–5 from the network (audit, academia, fintech) on equity tokens

Anti key-person by design: documented method, open-source core, advisory board and standards committee — the institution is built to outlive its founder.

THE ASK

€1.2M for 24 months — to close the proof, not to scale story

Enough capital to prove paid pull, conversion, cost-to-serve and channel — not enough to hide mistakes.

Engineering 55%	Founder/PM 13%	GTM / success 9%	Infra+security 8%	Trust+legal 7%	Contingency 8%
-----------------	----------------	------------------	-------------------	----------------	----------------

Totals 100% · SFP / SRL equity · semi-annual tranches against G2–G3 · lead engineer as condition precedent · 4-year founder vesting · ESOP 10–15% · non-dilutive grants as leverage

Investor

ticket €25–50k+ · cap €3–4M · data room open

Channel / association

white-label LOI · pilot cohort · standards seat

Regulated client

paid assessment €12–20k · reference

Public / partner

grants, public pilots · later-stage optionality

If the gates pass, this becomes a platform. If they don't, it stays a technical, standardizable asset with contained downside. That is the investment.

Gianluca Conte · contegianluca@hotmail.com · +39 340 1512072 · linkedin.com/in/contegianluca · github.com/Luke883i/RLA-ECNN

The deck opens the room; the room makes the case

Tag: ● ready · ◐ in progress · ○ to populate in the first 90 days

Commercial

- CRM pipeline
- ◐ LOIs & pilot contracts
- buyer interviews
- pricing tests
- rejections + reasons

Delivery

- assessment template
- ◐ sample audit pack
- logged hr/project margin
- before/after per client

Technical

- proof-flow demo
- repository
- negative tests
- receipt examples
- known gaps / baseline

Governance & Finance

- ◐ ETS/OpCo scheme + IP split
- ◐ licences · trademark · CLA · Col
- P&L 24m · use of funds · runway
- hiring plan · low/base/upside

Honest by construction: much of the room is still to populate — and populating it IS the next-90-days plan. What is already ● is verifiable today.

Full P&L (base case) and microfoundation of the numbers

€M	A0	A1	A2	A3	A4	A5
ARR (EoY)	—	0.06	0.35	1.2	4.0	10.7
Recognized revenue	0.05	0.17	0.42	1.1	3.1	7.7
Gross margin	—	70%	75%	80%	83%	85%
Opex	0.05	0.65	0.85	2.0	3.1	5.3
EBITDA	≈0	-0.53	-0.54	-1.1	-0.5	+1.2
FTE (EoY)	1	5	7	12	20	34

Formulas (auditable)

$ARR = SME \times \text{€}6k + \text{regulated} \times \text{€}50k + \text{sovereign maint.}$

$A5 \approx 1,200 \times 6k + 50 \times 50k + 4 \times 250k \approx \text{€}10.7M$

$GM = 1 - (\text{inference/cloud/support/channel}) / \text{rev}$

$\text{Payback} = \text{net CAC} / \text{annual gross profit}$

Drivers & sensitivity

- Killer variable 1: SME cohort conversion (assumed 30–40%; <15% → pivot mid-market, flagged at G3)
- Killer variable 2: time-to-proof-close (>M24 → round insufficient, flagged at G2)
- Killer variable 3: senior-days per €1 recurring revenue (the capacity inflection)

Expansion markets, exit optionality, risk and sources — out of the core story by design

Expansion (post-SME, gated)

- Regulated EU enterprise (DORA / NIS2): ACV €25–100k
- Digital government trust (eIDAS 2.0 / EUDI): sovereign deployments €0.5–2M
- Model-maker & infra partnerships (e.g. Mistral, Anthropic, ION) — category examples, no relationship implied

Exit optionality (illustrative, not a promise)

- GRC suites buy a proof substrate to differentiate
- Model makers / infra buy an admissibility wrapper
- PE roll-up of compliance-tech + channel + recurring revenue
- Or: dividend / cashflow case if it stays a defensible niche

Risk	Early signal	Gate / mitigation
Technical	proof chain doesn't close	G2 stop · lead engineer · negative tests
Commercial	assessments don't sell at target price	5 paid by M12 · pilot channel · pricing test
Capacity	delivery stays founder-bound	templating roadmap · channel handoff · senior-day metric
Open-core / IP	commons cannibalizes OpCo	licence, trademark, CLA, CoI, independent board
Key person	founder too central	CTO, ESOP, advisory, documented method

Sources (orders of magnitude, to verify in DD): AI Act Reg. (EU) 2024/1689 · DORA Reg. (EU) 2022/2554 (applicable 17 Jan 2025) · NIS2 (18 sectors) · eIDAS 2.0 Reg. (EU) 2024/1183 · EC SME data · vendor pages · open-core benchmarks (Mozilla, WordPress, Linux Foundation, Eclipse). Internal corpus: Whitepaper · Technical Due Diligence · CIPM · ROA. No economic figure is a guarantee of return or compliance.