

A-OSP

Computational Semantics for Proof Sovereignty

(Investment Pitch & Data room Annexes)

A local Proof Operating Environment that computes when AI-assisted work may answer, abstain, escalate or expose epistemic debt.

AI creates output. A-OSP computes proof status.

It does not open the black box. It moves the trust boundary outside the model.

Proof sovereignty · computational semantics · repo-governed engineering witness



ACT I — WHY IT EXISTS

Why A-OSP exists.

Problem, category, founder fit and capital-at-risk.

AI creates output. A-OSP computes proof status.

The bottleneck is no longer generating text — it is knowing whether a claim can be institutionally used, under what scope, with which evidence, and with which residual debt.

AI OUTPUT

fast · fluent · cheap

INSTITUTIONAL PROOF

slow · fragmented · expensive

A-OSP

local · reconstructable · governable


Core function

Compute when to answer, abstain, escalate or expose epistemic debt.

Generation is abundant. Admissibility is not.

AI-assisted work cannot be reconstructed.

The bottleneck is no longer generation. It is provenance, scope, gaps, review and accountability.



WHAT CHANGED

AI makes output abundant — documents, policies and summaries arrive faster than verification.

WHAT STAYED HARD

Who said it? Based on what? What is missing? Who reviewed it? Can it be defended?

WHAT A-OSP TARGETS

The reconstruction layer: source → claim → gap → receipt → artifact → witness → governance.

The failure is not only hallucination. It is unreconstructable institutional work.

A-OSP is a local Proof Operating Environment.

Not a chatbot, not a generic GRC dashboard, not a RAG wrapper. It is the proof layer around AI-assisted regulated work.

A-OSP is

- local proof substrate
- text-first virtual epistemic filesystem
- bounded model-call orchestration
- proof-and-governance workstation
- semantic status engine

A-OSP is not

- chatbot
- generic document repository
- model memory
- autonomous compliance engine
- legal certification

RAG retrieves context. A-OSP governs what retrieved context may become.

Compliance is the wedge. Computational proof infrastructure is the category.

The spend already exists. AI just made it acute.

A-OSP does not create a new budget line. It redirects money institutions already spend on reconstruction.

WHO PAYS TODAY

Compliance, audit, cyber and governance already pay consultants, auditors and data rooms to reconstruct evidence.

WHY NOW

AI made institutional output abundant overnight; reconstruction and defensibility became acute and recurring.

WHAT A-OSP CHANGES

It turns that recurring reconstruction spend into reusable proof capital instead of dead documents.

Already-budgeted spend · multi-regime duplication · audit-evidence pain · board / OdV / customer pressure

Compliance, audit, cyber and governance already pay for reconstruction — A-OSP makes it reusable.

Built by someone who lives the proof problem.

The founder profile sits where AI fails inside institutions: evidence, audit, accountability, controls and governance.

FOUNDER ASSET	WHY IT MATTERS
IT Audit & Internal Audit	Audit trail, controls, evidence, responsibility and remediation are native problems.
231 / OdV governance	Organizational liability, supervision, minutes, gaps and accountability known from practice.
Cybersecurity / NIS2 / identity	Connects proof to resilience, provider exposure, perimeter and recovery.
Privacy / AI governance	Moves AI governance from principles to operational evidence.
Applied philosophy of AI	Turns ignorance, uncertainty and responsibility into system objects.

Investment interpretation

Founder profile explains category coherence; the round exists to make the institution independent from the founder.

A-OSP is not AI looking for governance. It is governance, audit and proof discipline rebuilt around AI.

~€10k / 12 months: self-funded AI Operations Lab.

Personally funded frontier-model orchestration used as research, architecture, stress-test and build environment — not casual chatbot spend.

Inputs

- ~€10k personally funded
- 12 months of frontier-model use
- multi-model orchestration (operator, evaluator, integrator)
- founder as builder + integrator

Research

compression

Architecture

stress test

Governance

design

Repo method

& proof logic

What it proves

- Founder leverage
- AI orchestration skill
- Category coherence
- Compressed intellectual capital

What it does not prove

- Market validation
- Production readiness
- Security certification
- Benchmarked cost reduction

The round converts founder coherence into institution.


ACT II — WHY THIS ARCHITECTURE EXISTS

Why this architecture exists.

Theory lineage, local computability, computational semantics and epistemic debt.

The theory explains the architecture. The product is the witness.

Theory provides design grammar; the webapp is the engineering witness.



Claim discipline

A-OSP is ECNN-inspired and local-first. It is not claimed as a formally complete ECNN. The webapp is the witness: VFS, typed stripes, EQL/MONL, receipts, witness and governance.

Theory is the design grammar that explains why the product is built this way.

Why A-OSP does not promise universal truth.

It computes local proof status under declared constraints — and states precisely what it does not claim.

RICE

Non-trivial semantic properties are undecidable in general.

GÖDEL

No formal system closes all truth internally.

A-OSP RESPONSE

Compute admissible proof status under declared horizons — not total truth.

ALLOWED	FORBIDDEN
local proof status under horizon	universal truth machine
admissible transmission between scopes	global semantic completeness
computed unknown / contradiction / debt	guaranteed correctness

Local proof status under declared constraints — not universal truth.

Answer, abstain, escalate, expose debt.

A-OSP computes a discrete epistemic status for a claim under a declared horizon.

Given

- local typed proof graph
- declared horizon H
- admissible transmissions
- receipts / readback / validation state
- epistemic matrix M

The system computes

- supported
- unsupported
- contradiction
- out-of-scope
- review-required
- blocking debt

Not total truth. Computable epistemic status under horizon.

When semantic distinctions are lost, debt appears.

Epistemic debt is the residual obligation left when meaning is transformed, collapsed, labelled, transferred or reified without enough provenance, scope or validation to support downstream use.

OPERATION / FAILURE	DEBT
source missing	trace debt
label too broad	label debt
domain transfer	scale debt
claim promoted too early	reification debt
validation absent	validation debt

Undecidability is the formal limit. Epistemic debt is the governance residue.

“I don’t know” is not a phrase. It is a computed state.

The difference between a chatbot and a proof environment is where the unknown lives.

CHATBOT	A-OSP EPISTEMIC HEAD
“I don’t know” as style	unknown as computed state
confidence score	admissibility under proof graph
provider answer	local proof-state evaluation
refusal	abstention with scope and debt
model memory	typed local graph

answer
when supported

abstain
when not supported

escalate
when governance is needed

expose debt
when meaning is unresolved

From output to proof object.

An output becomes institutionally usable only by crossing a proof chain — a sequence of permissions, not a formatting pipeline.



FORBIDDEN EQUIVALENCE	CORRECT INTERPRETATION
output = proof	output is only a candidate surface
log = receipt	receipt requires shape, scope, hash / readback
export = witness	export needs witness-backed proof state
UI green = proof	green is earned only by proof predicates
model memory = SOT	SOT is local typed stripes

A proof chain is a sequence of permissions, not a formatting pipeline.


ACT III — WHAT HAS BEEN BUILT

What has been built.

A founder-built, AI-orchestrated engineering witness with repo governance and proof seams.

Not random prompting: governed AI development.

AI may generate drafts. Only code, tests, CI, runtime behavior and canonical docs become authority.



Authority pipeline

candidate (AI-generated) → typed (contracts) → tested (unit/E2E) → CI-gated (checks) → runtime-observed (behavior) → canonical (authority).

Engineering outcome

A founder-scale, AI-native method for building an unusual web-based epistemic operating environment.

AI-generated does not mean authoritative. Authority is earned.

The repo governs AI-assisted development.

Self-applied: the repository contains protocols that orient AI agents across local work, global vision, risk and proof boundaries.

AGENTS.MD

binding entry point for AI agents

IKANT PROTOCOL

anti-injection · required reads · authority
precedence

BUSSOLA

past · present · future · vision

TRADE-OFF

local vs global · effort · debt · alternatives

ICP GATES

issue alignment · roadmap · dataflow ·
breaking changes

AUTHORITY HIERARCHY

code > tests/CI > config > docs > issues > AI
plans

This does not prove enterprise readiness. It proves a method of layered governance for 1FTE + AI development.

A-OSP is the engineering witness.

The repository demonstrates the intended stack, proof boundaries and hardening path. Status: AS-IS / HARDENING, not production-certified.

BROWSER-BASED EPISTEMIC OS

web OS-like app surfaces

VIRTUAL EPISTEMIC FILESYSTEM

.aosp.txt typed append-only stripes

PARSER-SERVICE

EQL / MONL read/write/query boundary

DIRECTUS / POSTGRESQL

cache / index, not truth

D1/EQL PROOF SEAM

receipt, readback, explicit scope

GOVERNANCE DOCS

authority hierarchy and proof scripts

CI/CD CONTROL PLANE

classified, budgeted, hardening

EU-FIRST STACK

target posture (see A10)


Interpretation

aosp1 is not a finished enterprise product. It is a concrete witness that the stack, abstractions, proof seams and repo-governance method can exist together.

A witness, not a certificate.

Enterprise webapp + epistemic OS-like abstractions.

A-OSP uses dominant web technologies, then adds proof-native OS-like abstractions over the stack.



Web stack = enterprise substrate. A-OSP = epistemic / proof abstraction layer.

Plain text is not a downgrade. It is the resilience layer.

A-OSP does not use TXT instead of enterprise infrastructure. It uses TXT as the canonical proof substrate.

TECHNICAL BIAS	A-OSP RESPONSE
“TXT does not scale”	Text is SOT / event-log; parser, cache and DB handle query and performance.
“You need a database”	The database exists — it is cache / index / RBAC, not institutional truth.
“TXT is fragile”	Plain text is portable, diffable, inspectable, versionable and legacy-resilient.
“Where is schema?”	Stripes are typed; schema lives in parser / contracts / registry.
“Enterprise?”	Enterprise means recovery, auditability, access control, observability, reproducibility.

DB-resilient, not DB-less. LLM-resilient, not LLM-free. If the database, model or vendor changes, institutional proof remains readable.

First proof seam: receipt, readback, explicit scope.

D1/EQL is not claimed as full enterprise proof. It is the first proof seam to prevent false-green.

D1 PROOF

D1Receipt PASS + readback_hash + explicit session scope

EQL PROOF

EQLQueryReceipt + queryHash + resultHash + explicit scope

False-green rule

A completion modal, raw query success, a filled questionnaire or a local fallback atom do not equal proof.

Minimum proof substrate before D2–D5, not full proof spine.

The system must not look green before it is proof-green.

The goal is not to hide incomplete proof. The goal is to prevent incomplete proof from looking complete.

STRUCTURAL D1 TEST

no D1 green without PASS receipt

REQUIRED CHECKS

D1, dataflow, ontology, a11y, proof

E2E DETERMINISM

no static skip / hard wait in critical journeys

SEED BASELINE

canonical .aosp.txt inputs hashed

CAVEAT

full UI proof run remains opt-in / hardening in parts

No green without proof predicates.

Not every check means the same thing.

A-OSP classifies controls before claiming control: blocking, evidence, advisory, nightly, manual and deprecated checks differ in meaning.

CLASS	MEANING
blocking	must pass before merge
evidence	produces artifacts, no verdict
advisory	visible, not blocking
nightly	scheduled monitoring
manual	dispatch-only
deprecated	historical / archive

AS-IS inventory

- 81 workflow files
- 50 active · 31 archived
- 8 blocking budget
- 21 evidence · 11 advisory
- 8 nightly · 2 manual

Investment task: classify, budget and consolidate — this is what the round funds, not more workflow noise.

Honest status

Rich but still hardening. The goal is not more workflow noise; it is classified, budgeted, progressively consolidated control.

A-OSP classifies controls before claiming control.

ACT IV — SOVEREIGN & INVESTIBLE

Why it is sovereign and investible.

Provider resilience, EU-first posture, market wedge, moat stack, gates and ask.

Model ≠ memory. Database ≠ truth. Export ≠ proof.

Provider-agnostic proof sovereignty: proof identity stays local even while models, databases and clouds remain useful.

MODEL

replaceable processor for bounded transformations

DATABASE

rebuildable cache · index / RBAC / acceleration

UI / EXPORT

projection and representation, not proof by itself

.AOSP.TXT SOT

local institutional proof substrate — human + machine readable

A-OSP is not anti-provider. It is anti-provider-capture. Providers may process; they must not own proof, memory or institutional truth.

local truth



bounded processing



receipt & readback



witness & governance

Resilient by role separation, not by denial.

A-OSP does not eliminate providers. It limits what providers are allowed to mean.

COMPONENT	USED FOR	NOT ALLOWED TO OWN
LLM	bounded transformation under FAE (focus·attention·execution)	memory, truth, proof state
Database	index, cache, RBAC, acceleration	canonical knowledge
UI	navigation, projection, review	proof by appearance
Export	sharing, reporting	witness without receipt
.aosp.txt	local proof substrate	performance engine alone

AS-IS caveat

OpenAI / Anthropic adapters exist today. Target: provider registry + EU-first routing + local / self-hosted adapters. Security: no certified claim without P1/P2 controls, tenant isolation, recovery and provider-exposure testing.

Sovereignty is architectural first, EU-first second.

A-OSP is sovereignty-aligned because proof identity is provider-independent. EU-first is a deployment and dependency policy, not the core proof claim.

1 · Proof remains local

Institutional truth lives in .aosp.txt, never owned by a provider.

AS-IS

2 · Providers are replaceable

Models, databases and clouds are bounded, swappable processors.

AS-IS

3 · EU-first routing

EU / self-hosted adapters as the target routing posture.

TARGET

4 · Exposure register controls risk

Every dependency tracked with payload class, sensitivity and validation.

TARGET

Provider-independence is the proof claim; EU-first is the deployment policy.

Companies already pay for reconstruction.

The cost exists today: consultants, audit, compliance, data rooms, remediation, board reports. Value often dies in documents.

TODAY	WITH A-OSP
audit archaeology	reconstructable proof trail
dead reports	reusable proof capital
evidence rebuilt multiple times	one evidence object, many lenses
AI output hard to defend	receipt / gap / review state
fragile green dashboards	computed proof status

A-OSP makes compliance spend accumulate instead of evaporate.

Integrated compliance: one evidence object, many lenses.

Compliance is where evidence, absence of evidence, review and accountability are already paid for.

Why compliance first

- already budgeted
- multi-regime duplication
- audit evidence pain
- board / OdV / customer pressure
- proof and review already matter

Lenses over the same evidence object

231 / OdV

GDPR

NIS2

DORA

ISO 27001

Health & Safety

AI governance

(DORA / ISO map to RegTech & mid-market; 231 / H&S / GDPR are the SME wedge.)


Prior work matters

The compliance wedge is not opportunistic. It follows prior integrated-compliance work on multi-regime cost, evidence reuse and overlapping obligations.

Compliance is the wedge. Computational proof infrastructure is the category.

Paid proof first. Standard later.

Start with concrete paid objects. Build recurring proof capital. Prepare standardization only after adoption.



Standardization remains gated by adoption, not launched as overhead.

Proof is adopted through trust networks.

A-OSP sells proof, not automation. It enters through intermediaries who already speak evidence, governance and risk.

ODV / 231

organizational liability

INTERNAL AUDIT

controls + evidence

PRIVACY / AI GOV

data + accountability

CYBER / NIS2

security + resilience

UNIVERSITY / RESEARCH

governance + public interest

**ADVISORS / ASSOCIATIONS /
REGTECH**

distribution + trust

Trust networks reduce cold-start risk but do not replace paid validation.

Enter through people who already speak evidence, governance and risk.

Revenue is built by mixed market units.

Y3 revenue cannot come from manual assessments alone. It requires mixed market units and recurring proof capital.

MARKET UNIT	ANNUAL TARGET
Association unit	€150–300k
Advisor / OdV unit	€100–300k
Mid-market workflow	€150–350k
RegTech adapter / pilot	€150–500k
Enterprise pilot + maintenance	€150–300k

Delivery capacity — the mother metric			
	Y1	Y2	Y3
h / assessment	100	55	35
assessments / FTE	~7	~13	~21
founder-dependency	high	<60%	<40%
Manual delivery must decline through templates, proof objects and co-delivery.			

Revenue logic

recognized revenue = diagnostics + proof packs + maintenance + association units + advisor/OdV units + mid-market workflows + RegTech pilots / adapters.

Category unit: ecosystem / conformance / enterprise-platform ARR — post-M36 only.

Unit counts are bounded by delivery capacity, not by ambition.

The moat is layered, not singular.

A stack of theory, semantics, architecture, proof, governance, sovereignty and commercial trust — not a single feature.

THEORETICAL

ROA / RLA / CRC / ECNN design grammar

SEMANTIC

answer / abstain / escalate / debt

ARCHITECTURE

local SOT, VFS, parser, EQL/MONL

PROOF

receipt, readback, scope, witness, false-green prevention

GOVERNANCE

AGENTS, iKant, Bussola, ICP, authority hierarchy

SOVEREIGNTY

provider-replaceable, EU-first target posture

COMMERCIAL

compliance wedge + trust networks + standards-bound governance

As generation commoditizes, value migrates to the proof layer. A-OSP is strategically unnatural for incumbents built around data gravity and proprietary memory; sovereign, scientific and public-interest demand is optionality earned by the substrate — not near-term revenue.

Seven layers, each independently hard to copy.

The moat is not certainty. It is disciplined visibility.

A proof company must prove its own limits first.

LAYER	STATUS
.aosp.txt SSOT / DB cache	implemented design / hardening
D1/EQL seam	first seam active / not full proof spine
iKant / AGENTS protocol	active repo contract
AI-powered governance	partial / advisory / fallback
CI topology	generated / classified / still complex
E2E determinism	critical journeys hardened; long tail staged
EU-first stack	target posture, not AS-IS default
enterprise readiness	gated target (P1/P2)

A proof company must prove its own limits first.

4-year formation plan. M36 activation gate.

The plan keeps the activation event gated, not promised.

Y1

OpCo, CTO, P0 / D1 / EQL closure, first paid diagnostics

Y2

Repeatable delivery, P1 controls, advisor / association channels

Y3

Market units, RegTech / mid-market pilots, founder-dependence

M36

Activation review: enterprise pilots, strategic round, standards trigger

Y4+

Scale and standards Category option (see next slide)

THE ASK

Proof & Institutional Formation

€6.0M

gate: CTO M3 · SRL · P0 · ≥5 paid diagnostics

Strategic / Series-A

**separate,
at M36**

enterprise pilots + scale; not in this ask

M36 is not profitability. It is activation review for enterprise pilots, strategic-round readiness and standard scaling. Y1/Y2 must convert founder-operated repo governance into team-operable process.

The formation case funds the right to test category capture — it does not underwrite it.

Formation case — underwritten

- Y3: €5.8M recognized revenue
- M36: activation review, not profitability
- Goal: prove repeatable delivery + channels + proof seam

Scale case — post-M36

- €15–30M ARR
- Requires: enterprise pilots, RegTech adapters, advisor/OdV channels, recurring proof maintenance

Category case — 4+ years

- €100M+ ARR, 4+ years
- Requires: A-OSP proof format read by third-party tools, auditors, advisors, regulated institutions
- Revenue shifts from services/packs to platform, adapters, conformance mark / validator / enterprise workflows

€100M+ ARR is a post-M36 category-capture target, not a Y4 operating plan.

Fund discipline, not faith.

Founder leverage is not validation. Enterprise claims require gates. Economic claims require benchmarks. Semantic claims are bounded by horizon.

Activation gates

- P0 proof seam · P1 controls · P2 tenant/cost/failover
- full proof-flow demo repeatable
- provider-exposure registry + EU-first adapter plan
- 20 paid assessments · 2 channels · 1 RegTech pilot
- founder-dependence < 40%
- repo governance reproducible by CTO/team
- hardening-deferred index reduced
- revenue bridge validated
- standard-adoption signal: ≥3 third-party integrations, advisor/auditor workflows or formal conformance pilots

ALLOWED	FORBIDDEN
computational semantics under horizon	universal truth machine
founder-built proof asset	market-validated company
D1/EQL first seam	end-to-end proof-grade platform
enterprise-readiness plan	production-certified product
benchmark hypothesis	proven cost reduction

Additional no-go

No scale if the CI blocking/advisory map is unmaintained, the full proof-flow demo is not repeatable, or the hardening-deferred index grows without owner.

No activation without gates. No green without proof. No truth without horizon.

Fund computational proof sovereignty, not another AI wrapper.

A-OSP is a founder-built, AI-orchestrated engineering witness for local computational semantics: theory-informed, repo-governed, provider-resilient, EU-sovereignty aligned, honest about hardening, and designed to turn AI-assisted work into reconstructable institutional evidence.

Models generate.

A-OSP computes whether the result may be institutionally used.

Institutions decide.

But only after proof status, scope, debt and governance are visible.

Appendix — diligence pack.

Financial bridge · benchmarks · use of funds · governance · risks · founder · repo evidence · theory · computational semantics · sovereign stack · tests · hardening debt

A1

Revenue + runway

A2

Benchmarks

A3

Use of funds

A4

Governance

A5

Risks

A6

Founder CV

A7

Repo map

A8

Theory

A9

Comp. semantics

A10

Sovereign stack

A11

Tests

A12

Debt

Revenue bridge & runway: how Y3 is built and funded.

Illustrative. Revenue must come from market units and recurring proof capital, not manual assessments alone.

Y3 SOURCE	ILLUSTRATIVE REVENUE
Association units · 6 × €225k	€1.35M
Advisor / OdV units · 5 × €200k	€1.00M
Mid-market workflows · 8 × €150k	€1.20M
RegTech pilots / adapters · 3 × €250k	€0.75M
Assessment / packs / maintenance	€1.50M
Total	€5.80M

Runway (€6.0M round)

	Y1	Y2	Y3
Recognized rev.	0.6	1.85	5.8
Gross margin	50%	64%	72%
Opex	2.3	3.05	4.7
Net burn	2.0	1.9	0.5
Cum. cash left	4.0	2.1	1.6

€ M. Net burn = opex – gross profit. €6.0M funds the plan through ~M30–M36 if revenue tracks; the extension / bridge absorbs revenue slippage.

Interpretation

Y3 (€5.8M) is below full break-even: at €4.7M opex and 72% GM, break-even revenue is ~€6.53M. M36 is activation review, not profit guarantee.

Formation case stops at €5.8M recognized revenue. ARR scale case is not underwritten here; it is unlocked only after M36 activation.

Benchmark pack: prove the economics.

Economic sovereignty is a benchmark program, not a declared result.

LLM COST / PROOF PACK

€ / pack

PROVIDER SUBSTITUTION DELTA

quality–cost

EU-FIRST VS FRONTIER DELTA

quality–cost

READBACK SUCCESS

%

UNSUPPORTED CLAIM DETECTION

%

FALSE-GREEN DETECTION

%

PROOF-FLOW COMPLETION RATE

%

CI MINUTES / WORKFLOW COST

€ / run

HUMAN REVIEW TIME

hours

ASSESSMENT HOURS

h / assessment

Benchmark cost by workflow, then claim.

Use of funds: reconciled €6.0M formation round.

The ask funds institutional formation: CTO, engineering, security, delivery, governance, GTM and contingency.

AREA	AMOUNT
Engineering core	€2.05M
Product / security / cloud	€0.78M
Legal / IP / governance	€0.62M
Delivery / customer success	€0.57M
GTM / channel	€0.52M
Admin / insurance	€0.39M
Standards / offices / welfare	€0.57M
Contingency	€0.50M
Total	€6.00M

Capital ladder → gates

Formation €6.0M

T0 OpCo+CTO+P0 · T1 first clients+D1/EQL · T2 P1 controls+channel

Strategic / Series-A

separate, released at the M36 activation gate

€6M is not the cost of an MVP. It is the cost of converting a founder-built proof asset into an enterprise-operable company — staged against gates, not spent as a big-bang burn.

Governance is part of the product.

A proof standard cannot be captive to the same vendor that monetizes the runtime.

OpCo captures value

- product IP
- managed runtime
- enterprise modules
- customer contracts
- SLAs and integrations

Standards vehicle protects trust

- open proof specification
- conformance mark, not legal certification
- reference validator / open core
- training, standards, grants and public trust

Foundation-bound by design. Foundation-heavy only after adoption.

The standards vehicle is the path from product revenue to category economics: conformance mark, validator, open spec, ecosystem trust.

Risk and stop rules.

The pitch should make risk governable, not invisible.

RISK	STOP / MITIGATION
Founder consulting trap	Founder-dependence < 40%; delivery kit; co-delivery channels.
Manual assessment economics	No scale if h / assessment does not fall by Y2 close.
EU sovereignty overclaim	AS-IS vs target dependency register; no slogan.
LLM provider lock-in	Provider registry + EU-first adapters.
Database capture	.aosp.txt SSOT + rebuild path.
Founder-only knowledge	CTO / team runbooks + onboarding.
Governance theater / CI noise	Promote advisory to blocking only when stable; classify, budget, prune.
LLM cost overclaim	Benchmark costs by workflow; no 'irrisorio' guarantee.

Founder CV profile: category coherence.

The founder profile explains why proof, governance, AI and infrastructure sovereignty converge in one category.

CLUSTER	RELEVANCE
Audit / controls	Proof, reconstruction, accountability, audit trail.
231 / OdV	Organizational liability, governance, evidence, review.
Cybersecurity / NIS2 / identity	Resilience, provider exposure, perimeter, recovery.
Privacy / AI governance	Operationalizes AI governance as evidence infrastructure.
Applied philosophy of AI	Turns uncertainty and ignorance into system design.
Math / economics / game theory	Incentive-aware design for vendors, auditors, regulators.

Founder as originator of coherence, not permanent delivery engine.

Repo evidence map.

Map major pitch claims to concrete evidence surfaces in the aosp1 repository and data room.

CLAIM	EVIDENCE
VFS text-first	docs/02-FILESYSTEM.md
DB cache, not truth	docs/CANONICAL_SOURCES.md
dataflow discipline	docs/AI_DATAFLOW_GUIDE.md
agent governance	AGENTS.md
D1/EQL seam	docs/convergence/C1_D1_EQL_PROOF_SEAM.md
CI topology / required checks	docs/generated/ci-topology.md · docs/CI_REQUIRED_CHECKS.md
E2E determinism	docs/E2E_DETERMINISM.md
LLM providers AS-IS	backend/parser-service/src/llm/types.ts
Docker self-host path	backend/docker-compose*.yaml
Governance / LOA	Legal Operating Architecture full memo

Theory lineage.

The theory is not the product. It is the design grammar that explains the product.

LAYER	ROLE
Rice / Gödel	limits of universal semantic decision
ROA	observer governance, controlled reification, epistemic debt
RLA / CRC	local computable closure under horizon
ECNN / ECU	bounded epistemic units; epistemic head
iKant	normative meta-control: answer, abstain, escalate, expose debt
A-OSP	local-first proof OS: VFS, typed stripes, EQL/MONL, receipts, witness

Claim limit

A-OSP is ECNN-inspired. It is not claimed as a formally complete ECNN or as a research proof of ROA.

Computational semantics deep dive.

The system does not compute absolute truth. It computes admissibility under declared proof conditions.

Input

- claim C
- horizon H
- proof graph G
- transmission rules T
- epistemic matrix M

Output

- $\text{supported}(C, H)$
- $\text{unsupported}(C, H)$
- $\text{contradiction}(C, H)$
- $\text{out_of_scope}(C, H)$
- $\text{review_required}(C, H)$
- $\text{blocking_debt}(C, H)$

Formal statement

If L_i and L_j are locally computable under horizon H , and T_{ij} is a total computable transmission between admissible representations, then the two-node reticulum is operationally computable under H . This does not imply global semantic completeness.

Admissibility under declared proof conditions — not absolute truth.

Sovereign stack register.

Explicit AS-IS vs target posture per dependency layer — the basis of the EU-first dependency policy.

LAYER	AS-IS	TARGET
LLM	OpenAI / Anthropic	Mistral / EU-hosted / self-hosted adapters
DB	PostgreSQL / Directus	self-hosted; DB as cache
Cache	Redis	self-hosted
Runtime	Docker / Traefik	EU cloud / on-prem / sovereign hosting
DevOps	GitHub / npm	EU mirror + SBOM + signing
OS / devices	web-first	Linux / Android-first where feasible
Productivity	generic	privacy-aligned suite

Testing / anti-regression map.

A-OSP treats tests as proof surfaces, not decoration.

REQUIRED CHECKS

D1 SSOT, golden D1, ontology, dataflow, a11y, proof

E2E DETERMINISM

no static skips / no hard waits in critical journeys

FLAKE BUDGET

per-project tolerance; zero for critical gates

SEED BASELINE

hash canonical .aosp.txt inputs before E2E

FALSE-GREEN STRUCTURAL TEST

no D1 green without PASS receipt

DATAFLOW CONTRACT GUARDS

structural protection of canonical sources

Caveat

Full UI proof runs remain opt-in / hardening in parts. Testing moat = explicit anti-regression semantics, not a claim that all proof flows are finished.

Hardening debt index.

A-OSP does not hide hardening debt. It classifies it: visibility does not equal closure.

FINDING	MEANING
30 merged PRs scanned	deferred / out-of-scope / follow-up items extracted
47 distinct deferred items	grouped into 8 clusters
42 in-repo work items	tracked as bounded hardening work
18 follow-ups	eligible for adoption / wiring waves
9 post-MVP defers	explicitly not current MVP
Status caveat	visibility does not equal closure

Trust interpretation

Maturity honesty is a trust asset: the moat is disciplined visibility and layered risk reduction by a 1FTE non-engineer founder amplified by frontier AI — not certainty.