

A-OSP Build Data Room

Snapshot 2026-06-23 - HEAD 191ff93 - 255 days

FACT_

->

INFER_

->

GAP_

How an AI-orchestrated build measured its own proof debt - and what funding closes.

Scope: development process only - PR, issue, CI, ADR, entropy. Not product certification.

Method: FACT measured - INFER deduced - GAP declared blind spot.

SOT discipline:

WP = architecture / claim-boundary source.

TechDD = implementation posture source.

DR = repository/process evidence snapshot.

DR may add repository-backed facts, but must label them FACT / INFER / GAP and keep product-certification claims out of scope.

How this plugs into INP

INP tells the investment story. This data room turns it into repo facts.

Investment Pitch (INP) claim	Data Room (DR) evidence	Investment reading
AI creates output; A-OSP computes proof status	repo computes its own build status	the thesis is dogfooded
proof sovereignty	local debt/proof objects, not model memory	trust boundary lives outside the agent
proof capital	issues, gates, receipts, findings become reusable objects	debt can become an asset
funding ask	convergence throughput is the bottleneck	capital closes the measured gap

The round does not buy the idea. It buys convergence capacity: proof-seam closure, reproducibility, and team-operable governance.

ACT I - BUILD SIGNATURE

The repo is the first witness.

Signature, volume and febrile state - read as self-application evidence, not as certification.

Instrumented fever, not hidden entropy.

The repo is not healthy. The important fact is different: its unhealthy surfaces are measured, classified, owned and tied to convergence work.

One-glance state

self-declared health: 61.57/100, critical [FACT]
stable cold surfaces: 0 [FACT]
AI-authored code: 70.6% [FACT]
orphan debt on critical: 0 [FACT]
backlog: 46% explicit convergence, about 75% est.
[INFER]

Source discipline:

61.57 / redline / critical / hot = generated entropy report.

Contributor / PR / issue counts = GitHub snapshot at 191ff93.

Backlog convergence % = INFER, not measured FACT.

How to read it

high fever - but diagnosis and owner per symptom
repo consolidates itself; it does not grow features
governance measures; it does not cure
limit is convergence throughput, not awareness

The value is not optimism. It is measured proof debt with owner, temperature and gate.

Directed by one human, written by an agent.

The contributor signature explains both the strengths and the pathologies.

Contributions

Copilot SWE agent - 6,820 (70.6%)
Luke883i - 1,912 (19.8%)
github-actions bot - 914 (9.5%)
dependabot - 12 (0.1%)
one human merge-authority

Throughput

2,048 pull requests
1,679 merged (82%)
369 closed-non-merge
1,637 issues - 95% close-rate
81 issues open

Investor reading

- Risk: one human is the convergence bottleneck.
- Leverage: one human produced an out-of-scale witness.
- Funding use: turn founder-operated governance into team-operable governance.

One fact explains the rest: agent throughput created mass; human review tried to compress it into governed architecture.

The repository did not merely grow. It was forced to develop an immune system under abnormal AI-assisted throughput.

ACT II - THE INSTRUMENT

Debt is computable.

Before the numbers: the repo has tools to name, weight and assign proof debt.

The repo measures its own debt.

A-OSP does not only accrue development debt. It has deterministic tools to name, weight and assign it.

Measuring lanes

entropy lane - temperature in 5 classes
declared weighted formula for drift
self-audit that scores its own PRs
canonical taxonomy - 10 WTF-kinds
every finding bound to an owner issue

Product thesis mirrored in build

- Product: claim -> receipt -> owner -> no false-green.
- Build: finding -> ownerIssue -> gate -> convergence lane.

The proof discipline is not only described; it is applied to construction.

Instrument and thesis coincide: the product discipline is applied to the build process itself.

A self-declared temperature, not an external rating.

The temperature is regenerated in CI. Its value is reproducibility, not third-party authority.

Five classes

cold 0-20 - stable

warm 21-40 - monitor

hot 41-60 - convergence required

critical 61-80 - blocks proof on critical-path

redline 81-100 - quarantine / stop-the-line

Current reading

mean non-cold: 61.57 - critical

redline 2 - critical 12 - hot 10

warm 4 - cold 0

dominant debt: proof x11, observability x11

Do not overread the number as certification. Read it as a thermometer the repo can run again.

The weight formula is a confession of priorities.

The system fears false proof more than slow delivery. That is the A-OSP culture in one formula.

Weighted formula

proof_drift *20 - the maximum
runtime_drift *15 - authority drift *15
route-legacy-import-ui-issue *10
generated-doc *5 - feature-flag *5
temperature is weighted sum

What the formula says

Fear #1: false-green proof.
Fear #2: publication / divergence, partial runtime.
The two dominant anti-patterns are proof drift and observability drift.
Priorities are explicit, not hidden.

The highest weight is proof drift: a slow product is bad; a falsely green proof product is worse.

Every finding has an owner.

The hottest surfaces are not anonymous debt. They are assigned, sorted and bound to closure lanes.

Surface	Temp	Owner	Proof implication
MONL consumer/routes	98 redline	#3484	proof-path facade risk
DataIntake app	78 critical	#3462	D1 hardening
EQL query	78 critical	#3427 / #3430	retrieval proof seam
Artifact export	78 critical	#3472	export-as-proof risk
pipeline orchestrator	78 critical	#3471	progress-as-proof risk

The dominant debt is on the proof layer - exactly where the product addresses its market.

ACT III - THE HONEST GAP

Spec runs. Runtime lags.

The design ceiling is ahead of the witnessed floor. That is the investible gap.

The designed ceiling is ahead of the witnessed floor.

The question is not whether the architecture exists. It is how fast the floor converges with the ceiling.

Witnessed in runtime

- entropy diagnostics generated and owner-linked
- D1 proof route mounted
- D1ProofAppendService produces receipt/readback semantics
- d1_receipt_persistence evidenced
- frontend proofAppendD1 client wired
- DataIntake consumer path evidenced
- C1 false-green membrane generated
- D1 attack-vector tests present

Designed, not yet implemented

EQL is HARDENING / OPEN:
shared boundary and scope refusal exist,
but EQLQueryReceipt, hashes, proof scripts, CI gates and
proof UX remain open

D2-D5 receipts / PipelineReceipt remain open
ArtifactWitness remains open
GovernanceRun / Proof Spine remain open
CI/CD topology normalization remains open
MONL-TX remains redline / convergence target

The floor has moved: D1/C1 are no longer only designed; EQL and downstream proof composition are the closure frontier.

Funding closes the gap between coherent design and runtime proof, not a gap between no idea and idea.

The pathology is canonical; the ownership is unusual.

The dangerous version would be the same pathology without taxonomy, owner or gate.

Six anti-pattern families

- A1 duplication / over-generation
- A2 infinite-closure code - structural
- A3 dependency debt - uncovered
- A4 process hypertrophy - recognized
- A5 ADR drift / meta-overhead
- A6 route/spec drift

The fractal of AI

- AI-authored code has AI-style code-level debt.
- Provider/model/version changes are rarely measured.
- Human-in-architecture control is present but not yet enough.
- One reviewer compresses generated mass.

This is not brand damage. It is a controlled vocabulary for the build disease.

False-green is named before it is sold.

The repo names the failure A-OSP exists to prevent in customer work.

Redline and critical kinds

export-as-proof-copy - redline
progress-as-proof - redline
duplicate-canonical-claim - critical
parallel-runtime-path - critical
zombie-import / generated-doc-edited / fallback-hidden - critical

Why it matters

UI green cannot be proof.
Export cannot be witness.
Progress cannot stand in for closure.
A-OSP must forbid the false equivalences it found in its own build.

The repo discovered the product risk internally: progress, export or UI state masquerading as proof.

Review compresses generated mass.

A targeted sample shows the human compression effect of review over agent-generated code.

Sample PRs

#3522: 47k lines born; -89.5% in review
#3234: older rename; about 1.15 del/add
#3556: entropy PR; over 80% additions
#3623: icon cleanup

Funding implication

The founder work is not simply writing code.
It is compressing generated mass into governed architecture.
Funding buys parallel compression capacity.

Only four samples are measured, but the visible pattern is exactly the operating bottleneck: generated mass needs human convergence.

ACT IV - IMMUNE SYSTEM

From advisory thermometer to proof gate.

The immune system exists. The maturity gap is enforcement.

Out-of-scale immune system, staged enforcement.

A-OSP already has a build immune system. The roadmap is soft gates to hard gates to third-party readability.

Observe
entropy

->

Name
taxonomy

->

Own
issue

->

Gate
CI

->

Export
readable

Lanes and contracts

entropy convergence lanes
AI Agent Pattern Contract
PR-template-as-DoD
branch protection + required checks

Current maturity

D1/C1 membrane: generated + tested
within declared scope
Entropy / governance lanes: diagnostic
and partly advisory
CI/CD topology: instrumented, still
normalizing
EQL / D2-D5 / ArtifactWitness: open
convergence surfaces


Target maturity

EQL proof closure in CIdownstream
proof-consumer gates verified
D2-D5 receipts
ArtifactWitnessProof
Spinezero unknown-owner
governance surfaces
external reproducibility.

The immune system exists; capital accelerates enforcement.

Hallucination becomes debt; debt becomes guard.

A-OSP QA is not a promise to avoid every error. It is a discipline for converting errors into owned proof memory.



The loop

hallucination / drift / false-green -> owned debt ->
taxonomy -> negative test -> guard -> proof capital

Value proposition

Every closed false-green, every receipt seam and every
negative test converts epistemic debt into proof capital.

A-OSP gains value from the forces that make prompt-first AI fragile.

ACT V - INVESTMENT READING

Capital closes convergence.

Measured build debt into funding logic, governance architecture and stakeholder routing.

Canonical AI-authored pathology.

The repo is not weird because it is AI-authored. It is early: it shows what institutions will face as AI-assisted work scales.

Benchmark vs aosp1

code churn high and rising
clone / churn detection incomplete
issues-per-PR matched
tech-debt elevated
right comparison: governed vs unguided agent
throughput

The distinguishing move

A-OSP measures its own pathology.
The market typically discovers it after deployment.
The question is whether the immune system keeps pace.

Today it measures. Tomorrow it must block.

The limit is convergence throughput.

Every metric points to the same constraint: one human merge authority cannot converge AI-generated mass fast enough.

Constraint	Current signal	Funding response
one human merge authority	single reviewer / merger	convergence lead + reviewer cell
spec ahead of runtime	designed ceiling > witnessed floor	closure sprint
partial hard-gate maturity	D1/C1 evidenced, but EQL/downstream/CI topology not fully hard-gated	staged hardening: EQL CI, downstream gates, topology normalization
EQL / D2-D5 open	proof seam incomplete	P0/P1 proof-seam team
single snapshot	no derivative yet	recurring snapshot cadence

The fundable bottleneck is not awareness. It is convergence capacity.

Capital closes convergence, not invention.

A 90-120 day closure sprint should convert measured debt into runtime proof capital.

Current evidence	Funded workstream	Target
entropy 61.57 critical	entropy closure sprint	warm/cold mean
redline 2	MONL / proof-path convergence	redline 0
C1 generated false-green evidence	preserve + extend false-green membrane	regenerated C1, bypass 0
EQL HARDENING / OPEN	proof seam team	EQLQueryReceipt + hashes + proof CI
D2-D5 / ArtifactWitness open	downstream proof-chain team	PhaseReceipt + ArtifactWitness
one reviewer	convergence cell	founder-dependence reduced
4 churn samples	repo-wide analytics	aggregate churn median
CI/CD topology debt	governance topology cleanup	zero unknown-owner surfaces

Funding buys convergence capacity: EQL closure, downstream proof composition, governance topology normalization and team-operable review.

Build governance today; standards governance tomorrow.

The build process is the micro-model of the future proof commons.

DR proves now	LOA scales later
findings are inspectable	specification is foundation-held
proof debt has an owner	certification mark and governance can assign authority
false-green is prohibited	neutral standard cannot be vendor-captured
repo discipline is reproducible	third parties can read/write proof objects
founder governance is visible	OpCo/Foundation boundary formalizes control

No proof without inspectability; no standard without non-capture; no product margin without separation of capital.

Read the DR by stakeholder.

Each reader should extract a different decision from the same evidence.

Reader	What to read in DR	Next document
Investor	convergence gap is fundable	INP
CTO / engineering	proof seams are real but bounded	TechDD
Counsel / governance	traceability and owner discipline	LOA
Foundation / public capital	standard-readiness posture	LOA + INP
ROA reader	self-application witness	ROA

The DR should not be the first document. It should be the evidence appendix after INP.


Not healthy. Not failed. Governable.

The correct reading is governed acceleration.

A-OSP's first proof object is not a polished product demo.

It is an AI-assisted build that measured its own entropy, named its false-green risks, assigned debt to owners, and exposed the exact convergence work funding must close.

That is the investment reading: not optimism, not pessimism - governed acceleration.



The fever is the dataset. The traceability is the product. The convergence is the use of funds.